

# Class 5: Unsupervised learning and dimension reduction

Andrew Parnell  
`andrew.parnell@mu.ie`



PRESS RECORD

[https://andrewcparnell.github.io/intermediate\\_ML](https://andrewcparnell.github.io/intermediate_ML)

# Learning outcomes

Our first foray into **unsupervised learning**


- ▶ Reminder on some basic methods of dimension reduction
- ▶ An introduction to autoencoders and how they are used
- ▶ Some examples to show how they are used

## Set-up

- ▶ Suppose we have a high dimensional data set that we can write as a matrix  $X$  with  $n$  rows and  $p$  columns
- ▶ For example you might have a set of songs in the rows and details about them in the columns (number of streams, length, chord structure, etc)
- ▶ Plotting all of these data gets really hard as it's so high dimensional
- ▶ Can we extract some kind of meaning from the variability in these data? Are some columns linked? Are some of the row very different from the others
- ▶ Idea: approximate the matrix  $X$  with  $\tilde{X}$  which has far fewer columns but contains most of the information in  $X$


# Principal Components Analysis

- ▶ The oldest, simplest and most useful version of dimension reduction
- ▶ The method plots the data in high-dimensional space (mathematically) and then re-aligns the axes to match the spread of the data, so that the first axis has most of the variability, the second axes the second most, etc




## Example: the mnist data


Label: 5




Label: 4




Label: 0



Label: 1



# Running PCA with 2 dimensions



## Some notes on the previous slide

- ▶ The PCA never saw the digit number, it only saw the images (I used the digit numbers just for the final plot)
- ▶ I chose (and it's very common to choose) 2 principal components to plot, but you could choose many more
- ▶ The algorithm automatically orders the PCs and you can see how much variation each one captured in the data

```
summary(pca)$importance[,1:5] * 100
```

##	PC1	PC2	PC3	PC4	PC5
## Standard deviation	636.2889	540.7337	517.7845	454.7757	425.138
## Proportion of Variance	6.0790	4.3900	4.0260	3.1050	2.714
## Cumulative Proportion	6.0790	10.4690	14.4950	17.6000	20.314


- ▶ The command to run the PCA is really simple:

```
pca <- prcomp(x_train, center = TRUE, scale. = TRUE)
```

## Using PCA values as classifiers

- ▶ The first principal component is pretty good at classifying some of the digits (especially zero and one):

```
boxplot(projected_df$PC1 ~ projected_df$digit,  
        ylab = "PC1", xlab = "Digit")
```






## A different way of doing dimension reduction - autoencoders

- ▶ A very strange idea: why not fit a standard neural network model but have the features and the target to be the same!?
- ▶ How can this work? Rather than adding lots of extra weights in large hidden layers inside the model reduce the number of weights down to something much smaller
- ▶ Use those reduced weights as the dimension reduced version of the full data

# Structure of an Autoencoder

- ▶ An autoencoder consists of an encoder, a code layer, and a decoder
- ▶ The encoder compresses the input into a latent-space representation
- ▶ The decoder reconstructs the input data from the latent space



# Training an Autoencoder

- ▶ Because the model is really just a neural network, just like all the others we have fitted, we can use `keras`
- ▶ We structure it slightly differently because we want access to the latent space in the middle of the model
- ▶ Once fitted, we access the latent space in the middle and plot, or use it for whatever purpose we have

## Encoding Layer in R

- ▶ The encoding layer reduces the dimensionality of the input data
- ▶ We require an activation function, but we're not restricted to just one layer, we can use more if we want:

```
autoencoder %>%  
  layer_dense(units = 64, activation = 'relu',  
              input_shape = 784)
```

## Code Layer in R

- ▶ The code layer represents the compressed form of the input
- ▶ We use fewer units depending on how big a set of 'components' we want

```
autoencoder %>%  
  layer_dense(units = 32, activation = 'relu')
```

## Decoding Layer in R


- ▶ The decoding layer tries to reconstruct the original data from the code layer
- ▶ It usually mirrors the structure of the encoding layer

```
autoencoder %>%  
  layer_dense(units = 64, activation = 'relu') %>%  
  layer_dense(units = 784, activation = 'sigmoid')
```

## Running the model

- ▶ We then compile and fit the model the same as always
- ▶ But then we extract the weights of that middle layer and predict (a forward pass) those values for each of the observations (in the training or the test set)
- ▶ We can then plot them and see how the values change across the observations

# Example plot







## A reconstructed image

## 313/313 - 0s - 261ms/epoch - 835us/step

**Original image. Label: 7**



**Reconstructed image. Label: 7**




## What do we use an encoder for?

- ▶ You can use it to store a much smaller version of your data that contains most of the information
- ▶ You can use it to find observations that lie on the edge of the space and so might be considered anomalies
- ▶ If you find that some dimensions are capturing certain components of the data you can remove them (e.g. de-noising images)
- ▶ You could re-use the lower dimensional set as the features in another model which might be more efficient
- ▶ It can also be used to generate new data

# Variational autoencoders

- ▶ A cool extension of autoencoders are **variational encoders** (VAEs)
- ▶ VAEs extend AEs by introducing a probabilistic framework, where the encoder models the input data as a probability distribution in the latent space
- ▶ Usually a normal distribution is used with an estimated mean and variance matrix
- ▶ The fitting and the code gets a bit more complicated
- ▶ But the cool thing is that you can then use them to **generate** new training data

## Using VAEs to generate new images



## PCA vs AE vs VAE

- ▶ PCA is fast but might not capture complex structures as effectively
- ▶ Autoencoders are more flexible and can capture non-linear relationships
- ▶ Variational AEs are richer again and allow for the generation of new data
- ▶ All methods are available in R for dimension reduction

## Summary

- ▶ Dimension reduction very useful for all sorts of reasons, but especially if you have a lack of labelled data for supervised learning
- ▶ Pick a method that works for your data; I would always start with PCA
- ▶ We can use some of these methods later on for anomaly detection